

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.

COMPx270: Randomised and Advanced Algorithms


Lecture 5: Graph algorithms

Clément Canonne

School of Computer Science




THE UNIVERSITY OF
SYDNEY



Global Minimum Cut

Input: A connected, undirected graph $G = (V, E)$.

For a set $A \subseteq V$ let $\delta(A) = \{(u,v) \in E : u \in A, v \in V \setminus A\}$.




Aim: Find a cut (A, B) minimizing $|\delta(A)|$.

Global Minimum Cut

Input: A connected, undirected graph $G = (V, E)$.

For a set $A \subseteq V$ let $\delta(A) = \{(u,v) \in E : u \in A, v \in V \setminus A\}$.



Aim: Find a cut (A, B) of minimum cardinality.

Global Minimum Cut

Applications: Partitioning items in a database, identifying clusters of related documents, network reliability, network design, circuit design, TSP solvers.

Network flow solution.

- Replace every edge (u, v) with **two** directed edges (u, v) and (v, u) .
- Pick some vertex s and compute min s - v cut separating s from each other vertex $v \in V$.


Running time: $O((n-1) \cdot \text{MaxFlows})$

Global Minimum Cut

Running time: $O((n-1) \cdot \text{MaxFlows})$

Karger's Contraction Algorithm

Definition. A **multigraph** is a graph that allows multiple edges between a pair of vertices.



Karger's Contraction Algorithm

Definition. A **multigraph** is a graph that allows multiple edges between a pair of vertices.




Karger's Contraction Algorithm

Let $G=(V,E)$ be a multigraph (without self-loops).

The **contraction** of an edge $e=(u,v)\in E$ gives $G\setminus e$

- Replace u and v by single new **super-node** w
- Replace all edges (u,x) or (v,x) with an edge (w,x)
- Remove self-loops to w




Karger's Contraction Algorithm


Require: multigraph $G = (V, E)$

- 1: **while** $|V| > 2$ **do**
 - 2: Pick an edge $e \in E$ uniformly at random
 - 3: Contract it, and let $G \leftarrow G/e$
 - 4: **return** the cut defined by the remaining two vertices.
-


Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm




Karger's Contraction algorithm



Karger's Contraction algorithm




Karger's Contraction algorithm



The End.


Karger's Contraction Algorithm

Observation: An edge (u,v) contraction preserves the cuts (A,B) where u and v are both in A or both in B .



Karger's Contraction Algorithm

Observation: An edge (u,v) contraction preserves the cuts (A,B) where u and v are both in A or both in B .



Karger's Contraction Algorithm

Observation: An edge (u,v) contraction preserves the cuts (A,B) where u and v are both in A or both in B .

If $u,v \in A$ then $\delta_G(A) = \delta_{G \setminus e}(A)$.
(with u and v replaced with “ uv ”)

Karger's Contraction Algorithm

Observation: If (A,B) is a **minimum** cut, then we are less likely to choose an edge (u,v) crossing it!

Karger's Contraction Algorithm

Require: multigraph $G = (V, E)$

- 1: **while** $|V| > 2$ **do**
 - 2: Pick an edge $e \in E$ uniformly at random
 - 3: Contract it, and let $G \leftarrow G/e$
 - 4: **return** the cut defined by the remaining two vertices.
-

Claim. This algorithm has a **reasonable** chance of finding a min cut.

Key claim

Claim. If C is a min-cut, then the algorithm returns it with probability at least $2/n^2$.

Key claim

Claim. If C is a min-cut, then the algorithm returns it with probability at least $2/n^2$.

Proof.

Amplification

To amplify the probability of success, run the contraction algorithm many times.

Require: multigraph $G = (V, E)$, integer T

- 1: **for** $1 \leq t \leq T$ **do** \triangleright Use fresh (independent) random bits for each
- 2: Run Algorithm \square on G , let C_t be the output
- 3: **return** the smallest cut among all cuts C_1, \dots, C_T obtained

Amplification

To amplify the probability of success, run the contraction algorithm many times.

Claim. If we repeat the contraction algorithm $r \left\lceil \frac{1}{\epsilon} \right\rceil$ times with independent random choices, the probability that all runs fail is at most $(1/e)^r$.

Karger's Contraction Algorithm

Require: multigraph $G = (V, E)$

- 1: **while** $|V| > 2$ **do**
 - 2: Pick an edge $e \in E$ uniformly at random
 - 3: Contract it, and let $G \leftarrow G/e$
 - 4: **return** the cut defined by the remaining two vertices.
-

Running time?

Karger's Contraction Algorithm

Require: multigraph $G = (V, E)$

- 1: **while** $|V| > 2$ **do**
 - 2: Pick an edge $e \in E$ uniformly at random
 - 3: Contract it, and let $G \leftarrow G/e$
 - 4: **return** the cut defined by the remaining two vertices.
-

Running time?


The algorithm is iterated $O(n^2 \log(1/\delta))$ times... total running time $O(n^4 \log(1/\delta))$.

Karger's Contraction Algorithm

Can we do better?


Karger's Contraction Algorithm

Can we do better?



Improved algorithm

Improvement. [Karger-Stein 1996]



- Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when $n/\sqrt{2}$ nodes remain.
- Run contraction algorithm until $n/\sqrt{2}$ nodes remain.
- Run contraction algorithm **twice** on resulting graph, and return **best of two** cuts.

Improved algorithm

```
1: procedure MODIFIEDKARGER( $G = (V, E), s$ )
2:   while  $|V| > s$  do
3:     Pick an edge  $e \in E$  uniformly at random
4:     Contract it, and let  $G \leftarrow G/e$ 
5:   return  $G$ 
6: procedure KARGERSTEIN( $G = (V, E)$ )
7:   if  $|V| \leq 6$  then
8:     return a minimum cut           ▷ Brute-force computation
9:   Set  $s \leftarrow \lceil n / \sqrt{2} + 1 \rceil$ 
10:  ▷ Contraction
11:     $G_1 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
12:     $G_2 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
13:  ▷ Recursion
14:     $C_1 \leftarrow \text{KARGERSTEIN}(G_1)$ 
15:     $C_2 \leftarrow \text{KARGERSTEIN}(G_2)$ 
16:  return the smallest cut among  $C_1, C_2$ 
```

Improved algorithm: Karger-Stein

```
1: procedure MODIFIEDKARGER( $G = (V, E), s$ )
2:   while  $|V| > s$  do
3:     Pick an edge  $e \in E$  uniformly at random
4:     Contract it, and let  $G \leftarrow G/e$ 
5:   return  $G$ 
6: procedure KARGERSTEIN( $G = (V, E)$ )
7:   if  $|V| \leq 6$  then
8:     return a minimum cut           ▷ Brute-force computation
9:   Set  $s \leftarrow \lceil n/\sqrt{2} + 1 \rceil$ 
10:  ▷ Contraction
11:     $G_1 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
12:     $G_2 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
13:  ▷ Recursion
14:     $C_1 \leftarrow \text{KARGERSTEIN}(G_1)$ 
15:     $C_2 \leftarrow \text{KARGERSTEIN}(G_2)$ 
16:  return the smallest cut among  $C_1, C_2$ 
```

Running time?

Improved algorithm: Karger-Stein

Running time?

```
1: procedure MODIFIEDKARGER( $G = (V, E), s$ )
2:   while  $|V| > s$  do
3:     Pick an edge  $e \in E$  uniformly at random
4:     Contract it, and let  $G \leftarrow G / e$ 
5:   return  $G$ 
6: procedure KARGERSTEIN( $G = (V, E)$ )
7:   if  $|V| \leq 6$  then
8:     return a minimum cut      ▷ Brute-force computation
9:   Set  $s \leftarrow \lceil n / \sqrt{2} + 1 \rceil$ 
10:  ▷ Contraction
11:     $G_1 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
12:     $G_2 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
13:  ▷ Recursion
14:     $C_1 \leftarrow \text{KARGERSTEIN}(G_1)$ 
15:     $C_2 \leftarrow \text{KARGERSTEIN}(G_2)$ 
16:  return the smallest cut among  $C_1, C_2$ 
```

Improved algorithm: Karger-Stein

```
1: procedure MODIFIEDKARGER( $G = (V, E), s$ )
2:   while  $|V| > s$  do
3:     Pick an edge  $e \in E$  uniformly at random
4:     Contract it, and let  $G \leftarrow G / e$ 
5:   return  $G$ 
6: procedure KARGERSTEIN( $G = (V, E)$ )
7:   if  $|V| \leq 6$  then
8:     return a minimum cut           ▷ Brute-force computation
9:   Set  $s \leftarrow \lceil n / \sqrt{2} + 1 \rceil$ 
10:  ▷ Contraction
11:     $G_1 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
12:     $G_2 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
13:  ▷ Recursion
14:     $C_1 \leftarrow \text{KARGERSTEIN}(G_1)$ 
15:     $C_2 \leftarrow \text{KARGERSTEIN}(G_2)$ 
16:  return the smallest cut among  $C_1, C_2$ 
```

Success probability?

Improved algorithm: Karger-Stein

Success
probability?

```
1: procedure MODIFIEDKARGER( $G = (V, E), s$ )
2:   while  $|V| > s$  do
3:     Pick an edge  $e \in E$  uniformly at random
4:     Contract it, and let  $G \leftarrow G / e$ 
5:   return  $G$ 
6: procedure KARGERSTEIN( $G = (V, E)$ )
7:   if  $|V| \leq 6$  then
8:     return a minimum cut           ▷ Brute-force computation
9:   Set  $s \leftarrow \lceil n / \sqrt{2} + 1 \rceil$ 
10:  ▷ Contraction
11:     $G_1 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
12:     $G_2 \leftarrow \text{MODIFIEDKARGER}(G, s)$ 
13:  ▷ Recursion
14:     $C_1 \leftarrow \text{KARGERSTEIN}(G_1)$ 
15:     $C_2 \leftarrow \text{KARGERSTEIN}(G_2)$ 
16:  return the smallest cut among  $C_1, C_2$ 
```

Improved algorithm: Karger-Stein

Theorem. [Karger-Stein 1996] The Karger-Stein algorithm runs in time $O(n^2 \log n)$ and returns a min cut with probability at least $\Omega(1/\log n)$.

Corollary. The “best-of-T” Karger-Stein algorithm runs in time $O(n^2 \log^2 n)$ and returns a min cut with probability at least 99%.

Improved algorithm: Karger-Stein

Theorem. [Karger-Stein 1996] The Karger-Stein algorithm runs in time $O(n^2 \log n)$ and returns a min cut with probability at least $\Omega(1/\log n)$.

Corollary. The “best-of-T” Karger-Stein algorithm runs in time $O(n^2 \log^2 n)$ and returns a min cut with probability at least 99%

Best known. [Karger 2000] $O(m \log^3 n)$.

And now, for something completely different

And now, for something completely different?


Theorem. An undirected graph $G=(V,E)$ has at most _____ distinct min cuts.

And now, for something completely different?

Theorem. An undirected graph $G=(V,E)$ has at most $n(n-1)/2$ distinct min cuts.

And now, for something completely different?

Theorem. An undirected graph $G=(V,E)$ has at most $n(n-1)/2$ distinct min cuts. And this is tight.



And now, for something completely different?

Theorem. An undirected graph $G=(V,E)$ has at most $n(n-1)/2$ distinct min cuts.

Proof.

And now, for something completely different?

How mind-blowing is this? Analyzing an algorithm proved a seemingly unrelated mathematical statement. 🧠

Compare this with last time, and the probabilistic method ("proof without algorithm!").

And now, for something else completely different!

Minimum Spanning Tree in (Expected) Linear Time: Karger-Klein-Tarjan.

