

Problems 1, 2 require you to have read the lecture notes or watched the lecture, but you should be able to attempt them on your own after that. Similar for 3 and 4.

Problem 5 takes some time, but is important to check you understand how the Morris counter works. The last subquestion is useful to make you think about how to implement some general building blocks requiring sampling bias bits, but a little bit harder.

Problem 6 is short, once you have the right idea. You can skip it if you don't see how to easily solve it, but read the solution later.

Problem 7 is quite involved, but useful to have seen (it introduces an interesting, "well known" algorithm): it is fine to skip it, but in that case still have a look at the algorithm (and the solution) afterwards.

Warm-up

Problem 1. Go through the "median-of-means" proof for the Morris counter, to prove the statement about the result.

Problem 2. Instead of using the "median-of-means" trick to boost the accuracy of the Morris counter, what happens if we were to do the opposite and use the "mean-of-medians"?

Problem 3. If we were to use truly random hash functions instead of a strongly universal hash family in the BKJST algorithm, what would the space complexity become?

Problem 4. Check your understanding: why are we using a hash function g in the BKJST algorithm? What would happen if we were to store j in the bucket B , instead of $g(j)$?

Problem solving

Problem 5. Analyse the "careful variant" of the Morris counter, where instead of doubling C with probability $1/C$, we multiply C by $1 + \alpha$ with probability $1/(\alpha C)$.

- Compute the expectation of C at the end of the algorithm.
- Compute its variance, and conclude by Chebyshev.
- Explain how you would set α to get a $(1 + \epsilon)$ -factor estimate of the true count with probability at least $1 - \delta$ using the median trick, and give the resulting space complexity, *almost* proving the theorem stated in the lecture (but with a multiplicative instead of additive $\log(1/\delta)$).

- d) Explain how you would set α to get a $(1 + \varepsilon)$ -factor estimate of the true count with probability at least $1 - \delta$ *without* using the median trick, and give the resulting space complexity, *actually* proving the theorem stated in the lecture (with an additive $\log(1/\delta)$).
- e) (\star) How would you actually implement the increment step (i.e., how, given random uniform bits, would you “multiply C by $1 + \alpha$ with probability $1/(\alpha C)$ ”)?

Problem 6. Given a stream σ of length m and $\varepsilon \in [0, 1]$, let

$$H_\varepsilon(\sigma) = \{j \in [n] : f_j \geq \varepsilon \cdot m\}$$

denote the set of ε -heavy hitters of σ . Modify the Misra–Gries algorithm to make it output a set $H \subseteq [n]$ such that $H_\varepsilon(\sigma) \subseteq H \subseteq H_{\varepsilon/2}(\sigma)$. (That is, the algorithm outputs every ε -heavy hitter, and everything it outputs is at least an $(\varepsilon/2)$ -heavy hitter.) Your algorithm should be one-pass, and use space $O(\log(mn)/\varepsilon)$.

Problem 7. Consider the following “Bottom- k ” algorithm for the Distinct Elements problem, where $k \geq 1$ is a parameter.

```

1: Pick a hash function  $h: [n] \rightarrow [0, 1]$  from a strongly universal hash family
   ▷ Technically, from  $h: [n] \rightarrow \{0, 1/N, 2/N, \dots, 1\}$  where  $N = \text{poly}(n)$  is large enough to not have to worry about collisions.
2: Set  $z = (1, 1, \dots, 1) \in \mathbb{R}^k$ 
3: for all  $1 \leq i \leq m$  do
4:   Get new element  $a_i \in [n]$  of the stream
5:    $z \leftarrow k$  smallest values among  $z_1, \dots, z_k, h(a_i)$ 
6: return  $\hat{d} \leftarrow \frac{k}{\max(z_1, \dots, z_k)}$ 

```

- a) Show that, for $k = \Theta(1/\varepsilon^2)$, the above algorithm returns $(1 \pm \varepsilon)$ -estimate of the number of distinct elements d , with probability at least 99%. (To do so, define X_i as the indicator that $h(a_i) < \frac{k}{(1+\varepsilon)d}$, and use Chebyshev on $\sum_i X_i$.)
- b) What is the space complexity of the algorithm?

Advanced

Problem 8. Given a stream σ of length m where each element a_i is a vector in $\{-1, 1\}^d$, our goal is to estimate how large the mean vector

$$\bar{a} := \frac{1}{m} \sum_{i=1}^m a_i \in [-1, 1]^d$$

is: that is, to obtain a multiplicative estimate of $\|\bar{a}\|_2$ within a factor 2, with probability at least 99%. Assuming that the algorithm does not “pay” the cost of storing the random bits it uses, describe an approach to do this with very small space complexity. What is the number of “free random bits” your algorithm uses?