

You can try Problems 1 to 3 at home after the lecture; if you are stuck with one, ask during the tutorial for guidance, but attempt it on your own afterwards.

Problem 4 is quite technical: good practice, but don't have time during the tutorial or afterwards, feel free to skip it and only read the solution.

Problems 5 and 6 (Advanced) are good practice, a bit longer and less guided. Discuss them during the tutorial: attempt to fill in the gaps on your own or in groups if you don't have time to finish during the tutorial.

Warm-up


Problem 1. Try various parameters for Theorem 60: plot, for $\beta \in (0, 1)$, the bound, when

- $C^* = 0, n = 1000$
- $C^* = 10, n = 1000$
- $C^* = 100, n = 1000$
- $C^* = 1000, n = 1000$

Do the same for Theorem 61.

Solution 1. Try it at home! Here in Mathematica:

```
In[3]:= Plot[{ (Log[1000]) / Log[2 / (1 + beta)],
  (10 * Log[1 / beta] + Log[1000]) / Log[2 / (1 + beta)],
  (100 * Log[1 / beta] + Log[1000]) / Log[2 / (1 + beta)],
  (1000 * Log[1 / beta] + Log[1000]) / Log[2 / (1 + beta)],
  (Log[1 / beta] + Log[1000]) / (1 - beta),
  (10 * Log[1 / beta] + Log[1000]) / (1 - beta),
  (100 * Log[1 / beta] + Log[1000]) / (1 - beta),
  (1000 * Log[1 / beta] + Log[1000]) / (1 - beta)},
{beta, 0, 1}, PlotLegends -> "Expressions"]
```



Problem 2. Assume you know both n and (an upper bound on) C^* in advance. How would you set β in the MWU? In the Randomised MWU?

Solution 2. Given explicit values: differentiate the expressions to find the minimum (or find the minimum numerically).

To find a reasonable approximation (up to a factor 2): choose β to balance the two terms in the numerator,

$$C^* \log(1/\beta) = \log n$$

This might not be the exact minimum, but will be within a constant factor, and is much simpler to derive.

Problem 3. Prove Fact 56.3: namely, consider $n = 2$ experts, one predicting always 0 and the other always 1, and consider all 2^T possible sequences (u_1, \dots, u_T) . Show that for any deterministic algorithm A , there exists a sequence on which the algorithm makes T mistakes, and use this to conclude.

Solution 3. Same reasoning and construction as Fact 56.1, but need to also show that $C^*(T)$ is at most $T/2$ (best expert will make at most $T/2$ mistakes). Note that since one expert always outputs 0 and the other always 1, the best expert will make at most $T/2$ mistakes: denote $S_1 = \{t \in [T] : u_t = 0\}$ and $S_2 = [T] \setminus S_1$, expert 1 will be correct $|S_1|$ times and expert 2 correct for $|S_2|$ times –

$$\max\{|S_1|, |S_2|\} = \max\{|S_1|, T - |S_1|\} \geq T/2$$

Suppose given algorithm A (you can predict what it will output every step)

- 1) observes $\{0, 1\}$, predict A 's output \hat{u}_1 . Set $u_1 \leftarrow 1 - \hat{u}_1$.
- 2) observes $\{u_1, 0, 1\}$, predict A 's output \hat{u}_2 . Set $u_2 \leftarrow 1 - \hat{u}_2$.
- 3) observes $\{u_1, u_2, 0, 1\}$, predict A 's output \hat{u}_3 . Set $u_3 \leftarrow 1 - \hat{u}_3$.
- etc. until T .

So your algorithm will make T mistakes while best expert will make at most $T/2$. The factor 2 is necessary for deterministic algorithms.

Problem solving

Problem 4. (*) Suppose $C^* = C^*(T)$ is known in advance. We will show how to modify the MWU algorithm to achieve

$$C(T) \leq 2C^* + O(\sqrt{C^*(T) \log n} + \log n)$$

- a) Argue that, if $C^* \leq \log n$, we are done.
- b) Suppose $C^* > \log n$. Show how to achieve the desired bound by setting $\beta = 1 - \varepsilon$, for some suitable $\varepsilon = \varepsilon(C^*, n)$.
- c) Conclude.

Solution 4.

a) If $C^* \leq 4 \log n$, then we set $\beta = \frac{1}{2}$.

$$\frac{C^* \log \frac{1}{\beta} + \log n}{\log \frac{2}{1+\beta}} \leq 2.41(C^* + \log n) \leq O(\log n).$$

b) We need some approximation facts first. For $\varepsilon \in (0, 0.5)$, we have $\varepsilon \leq \log \frac{1}{1-\varepsilon} \leq 2\varepsilon$ and

$$\frac{\log \frac{1}{1-\varepsilon}}{\log \frac{1}{1-\varepsilon/2}} = 2 + \frac{\varepsilon}{2} + O(\varepsilon^2) \leq 2 + \varepsilon.$$

by series expansion. We restrict $1 - \beta = \varepsilon < \frac{1}{2}$. We proceed with the calculation:

$$\begin{aligned} C^* \frac{\log \frac{1}{\beta}}{\log \frac{2}{1+\beta}} + \frac{\log n}{\log \frac{2}{1+\beta}} &= C^* \frac{\log \frac{1}{1-\varepsilon}}{\log \frac{1}{1-\varepsilon/2}} + \frac{\log n}{\log \frac{1}{1-\varepsilon/2}} \\ &\leq C^* \cdot (2 + \varepsilon) + \frac{\log n}{\varepsilon} \\ &= 2C^* + C^*\varepsilon + \frac{\log n}{\varepsilon} \\ &\leq 2C^* + C^* \sqrt{\log n / C^*} + \frac{\log n}{\sqrt{\log n / C^*}} \\ &= 2C^* + 2\sqrt{C^* \log n} \end{aligned}$$

Note that we need $\varepsilon = \sqrt{\log n / C^*} < \frac{1}{2}$, or, equivalently, $\frac{\log n}{C^*} < \frac{1}{4}$. (This motivates in hindsight question a)).

c) Combining the two, we have that

$$\frac{C^* \log \frac{1}{\beta} + \log n}{\log \frac{2}{1+\beta}} \leq 2C^* + O\left(\sqrt{C^* \log n} + \log n\right).$$

Problem 5. We again have n experts, each making a binary prediction at each time step. However, we would like to make sure we do well even if the best expert does badly overall, as long as for each “chunk” $I_{t_1, t_2} = \{t_1, t_1 + 1, \dots, t_2 - 1, t_2\}$, we do well compared to the best expert *for this chunk*.

To try and get this, consider the variant of the MWU, where we only penalise an expert by multiplying its weight by $1/2$ if its current weight is at least $1/3$ of the average weight of all experts.

We want to show that for every $1 \leq t_1 \leq t_2 \leq T$, the maximum number of mistakes $C(t_1, t_2)$ that the algorithm makes over I_{t_1, t_2} is at most $O(C^*(t_1, t_2) + \log n)$,

where $C^*(t_1, t_2)$ is the number of mistakes made by the best expert *in that chunk*. (Considering $\beta \in (0, 1)$ to be a constant, e.g., $\beta = 1/2$.)

- a) Write down the algorithm.
- b) Consider any chunk $I = I_{t_1, t_2}$, and let $t \in I$ be a time step where a mistake is made. Let W_t be the total weight at the beginning of step t , and W_G, W_B, W_L be the total weight of (1) experts who made a mistake, (2) experts who did not, and (3) experts who made a mistake but have weight less than $\frac{1}{3} \cdot \frac{W_t}{n}$. Bound the weight W_{t+1} at the end of step t as a function of W_G, W_B, W_L, β .
- c) Bound the weight W_{t+1} at the end of step t as a function of W_t, β : show that

$$W_{t+1} \leq \frac{5+\beta}{6} W_t$$

- d) Give a lower bound on the weight w_{i, t_1} of *any* expert i at time t_1 (start of the chunk). Namely, show that

$$w_{i, t_1} \geq \frac{\beta W_{t_1}}{3n}, \quad 1 \leq i \leq n$$

- e) Letting W_{t_1} the total weight at the beginning of the chunk, and W_{t_2} at the end, show that

$$W_{t_2} \geq \beta^{C^*(t_1, t_2)} \cdot \frac{\beta W_{t_1}}{3n}$$

- f) Conclude.

Solution 5.

- b), c) If we make a mistake at time t , the algorithm will have $W_G > W_B$ and thus

$$W_G \geq \frac{1}{2}(W_G + W_B) = \frac{1}{2} W_t.$$

And we have $W_L \leq \frac{1}{3} W_t$,

$$\begin{aligned} W_{t+1} &= \beta W_G + W_B + (1 - \beta) W_L \\ &= W_G + W_B + (1 - \beta)(W_L - W_G) \\ &\leq W_t + (1 - \beta) \left(\frac{1}{3} W_t - \frac{1}{2} W_t \right) \\ &= W_t - \frac{1 - \beta}{6} W_t = \frac{5 + \beta}{6} W_t. \end{aligned}$$

- d) Denote by \tilde{W} the total weight when expert i got penalised before t_1 . Since we only decrease weights as time progress, we have $\tilde{W} \geq W_{t_1}$. If i gets penalised and denote its weight at that time \tilde{w}_i , it must at the time have

$$\tilde{w}_i \geq \frac{\tilde{W}}{3n}.$$

And $w_{i,t_1} = \beta \tilde{w}_i \geq \frac{\beta \tilde{W}}{3n} \geq \frac{\beta W_{t_1}}{3n}$. If expert i has not been penalised up until t_1 , we know that $w_{t_1} = 1$ and therefore $w_{t_1} = 1 \geq \frac{\beta W_{t_1}}{3n}$.

- e) The best expert makes at most $C^{*(t_1, t_2)}$ mistakes; let j be the index of that expert. Then,

$$w_{j,t_2} \geq \beta^{C^{*(t_1, t_2)}} w_{j,t_1} \geq \beta^{C^{*(t_1, t_2)}} \frac{\beta W_{t_1}}{3n}.$$

Finally,

$$W_{t_2} = \sum_{i=1}^n w_{i,t_2} \geq w_{j,t_2} \geq \beta^{C^{*(t_1, t_2)}} \frac{\beta W_{t_1}}{3n}$$

Advanced

Problem 6. In the setting of the MWU, we have n experts, each making a binary prediction at each time step. Now, assume that we know that, for every $1 \leq k \leq n$, the k -th expert makes at most k mistakes.

- What bound can you show on $C(T)$ when running the MWU algorithm with parameter β ?
- What bound can you show on $\mathbb{E}[C(T)]$ when running the Randomised MWU algorithm with parameter β ?

Solution 6.

- We can use an analysis very similar to that given in class for the MWU algorithm. On the one hand, if the algorithm makes C mistakes then after these mistakes the total weight W of all experts will be at most $n \left(\frac{1+\beta}{2} \right)^C$. On the other hand, we now know that the k -th expert makes at most k mistakes, so the lower bound on the total weight we have is $W \geq \beta + \beta^2 + \dots + \beta^n = \beta \cdot \frac{1-\beta^n}{1-\beta}$. Solving the inequality

$$\beta \cdot \frac{1-\beta^n}{1-\beta} \leq n \left(\frac{1+\beta}{2} \right)^C$$

we obtain

$$C(T) \leq \frac{\log_2 \frac{1-\beta}{\beta(1-\beta^n)} + \log_2 n}{\log_2 \frac{2}{1+\beta}} = \frac{\ln \frac{1-\beta}{\beta(1-\beta^n)} + \ln n}{\ln \frac{2}{1+\beta}}$$

as our bound on the number of mistakes.

- b) The analysis is analogous to that of the Randomized MWU algorithm from the lecture. Let F_i denote the fraction of weight at the i -th trial on experts giving an incorrect advice, so that $C = \sum_{i=1}^T F_i$. On the one hand, we have that W (the final total weight of all experts) equals $n \prod_{i=1}^T (1 - (1 - \beta)F_i)$. On the other hand, we know that expert k has weight at least β^k , so here again $W \geq \beta + \beta^2 + \dots + \beta^n = \beta \cdot \frac{1-\beta^n}{1-\beta}$. Putting these together as in the lecture,

$$\mathbb{E}[C(T)] = \sum_{i=1}^T F_i \leq \frac{\ln \frac{1-\beta}{\beta(1-\beta^n)} + \ln n}{1 - \beta}.$$